

```
# -----
# 1) Imports
# -----
import os
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error, r2_score

# SHAP (install if missing)
try:
    import shap
except ImportError:
    !pip -q install shap
    import shap

# Colab upload helper (safe if not in Colab)
try:
    from google.colab import files
    IN_COLAB = True
except ImportError:
    IN_COLAB = False

# -----
# 2) Load Dataset (path or upload)
# -----
CANDIDATE_PATHS = [
    "data/salary_dataset.csv",
    "Salary_Data.csv",
    "salary.csv"
]

df = None
for p in CANDIDATE_PATHS:
    if os.path.exists(p):
        df = pd.read_csv(p)
        print(f"Loaded dataset from: {p}")
        break

if df is None:
    if IN_COLAB:
        print("File not found at common paths. Please upload your CSV.")
        uploaded = files.upload()
        fname = next(iter(uploaded))
        df = pd.read_csv(fname)
        print(f"Loaded dataset from upload: {fname}")
    else:
        raise FileNotFoundError("Dataset not found. Update CANDIDATE_PATHS or provide a valid path.")

print(df.head())

# -----
# 3) Normalize headers & clean
# -----
# Strip whitespace first
df.columns = df.columns.str.strip()

# Robust renaming: handle both underscores and spaces variants
# We'll normalize to lowercase for matching, then rename original columns.
lower_cols = {c: c.lower() for c in df.columns}
rename_targets = {
    "education_level": "Education",
    "education level": "Education",
    "years_of_experience": "YearsExperience",
    "years of experience": "YearsExperience",
    "job_title": "JobTitle",
    "job title": "JobTitle",
}
rename_map = {}

-----
```

```
for original, low in lower_cols.items():
    if low in rename_targets:
        rename_map[original] = rename_targets[low]
# Apply renaming
df = df.rename(columns=rename_map)

# Required columns after renaming
required_cols = ['Age', 'Gender', 'Education', 'YearsExperience', 'Salary']
missing = [c for c in required_cols if c not in df.columns]
if missing:
    raise ValueError(f"Missing expected columns: {missing}\nCurrent columns: {list(df.columns)}")

# Basic cleaning
df = df.drop_duplicates().copy()

# Normalize string columns
if 'JobTitle' in df.columns:
    df.loc[:, 'JobTitle'] = df['JobTitle'].astype(str).str.strip().str.title()

df.loc[:, 'Education'] = (
    df['Education'].astype(str)
    .str.replace('.', '', regex=False)
    .str.strip()
    .str.title()
)
df.loc[:, 'Gender'] = df['Gender'].astype(str).str.strip().str.title()

# Drop rows with critical missing values
df = df.dropna(subset=['Salary', 'Age', 'Gender', 'Education', 'YearsExperience']).copy()

# -----
# 4) Encoding
# -----
# Gender label-encode
le_gender = LabelEncoder()
df.loc[:, 'Gender'] = le_gender.fit_transform(df['Gender'])

# Education ordinal map with fallback
edu_order = ['High School', 'Bachelor', 'Master', 'Phd']
edu_map = {name.title(): i+1 for i, name in enumerate(edu_order)}
df.loc[:, 'Education'] = df['Education'].map(edu_map)

if df['Education'].isna().any():
    # Assign a value just above the max defined ordinal for unknown categories
    max_known = max(edu_map.values()) if edu_map else 0
    df.loc[df['Education'].isna(), 'Education'] = max_known + 1

# -----
# 5) Quick EDA
# -----
print("\nDescriptive statistics:")
print(df.describe(include='all'))

plt.figure()
sns.histplot(df['Salary'], kde=True)
plt.title("Salary Distribution")
plt.xlabel("Salary")
plt.ylabel("Count")
plt.show()

plt.figure(figsize=(8,6))
corr = df[['Age', 'Gender', 'Education', 'YearsExperience', 'Salary']].corr(numeric_only=True)
sns.heatmap(corr, annot=True, cmap="coolwarm")
plt.title("Feature Correlation Heatmap")
plt.show()

# -----
# 6) Features, Split, Scale
# -----
features = ['Age', 'Gender', 'Education', 'YearsExperience']
X = df[features].copy()
y = df['Salary'].copy()

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)

numeric_cols = ['Age', 'YearsExperience']
```

```
scaler = StandardScaler()

X_train_scaled = X_train.copy()
X_test_scaled = X_test.copy()
X_train_scaled.loc[:, numeric_cols] = scaler.fit_transform(X_train[numeric_cols])
X_test_scaled.loc[:, numeric_cols] = scaler.transform(X_test[numeric_cols])

# -----
# 7) Train & Evaluate
# -----
models = {
    "Linear Regression": LinearRegression(),
    "Random Forest": RandomForestRegressor(n_estimators=200, random_state=42)
}

results = {}
fitted_models = {}

for name, model in models.items():
    model.fit(X_train_scaled, y_train)
    y_pred = model.predict(X_test_scaled)
    rmse = float(np.sqrt(mean_squared_error(y_test, y_pred)))
    r2 = float(r2_score(y_test, y_pred))
    results[name] = {"RMSE": rmse, "R^2": r2}
    fitted_models[name] = model

results_df = pd.DataFrame(results).T.sort_values(by="RMSE")
print("\nModel performance:")
print(results_df)

# -----
# 8) Feature Importance (RF)
# -----
rf = fitted_models["Random Forest"]
importances = pd.Series(rf.feature_importances_, index=features).sort_values(ascending=True)

plt.figure()
importances.plot(kind='barh', title='Feature Importance (Random Forest)')
plt.xlabel("Importance")
plt.ylabel("Feature")
plt.show()

# -----
# 9) SHAP Explainability (RF)
# -----
shap.initjs()
explainer = shap.TreeExplainer(rf)
shap_values = explainer.shap_values(X_test_scaled)

# Bar summary
shap.summary_plot(shap_values, X_test_scaled, feature_names=features, plot_type="bar")
# Beeswarm summary
shap.summary_plot(shap_values, X_test_scaled, feature_names=features)

# -----
# 10) Conclusion
# -----
best_model = results_df.index[0]
print(f"""
Summary:
- Headers normalized (handled 'Education Level' → Education, 'Years of Experience' → YearsExperience, 'Job Title' → JobTitle).
- Encoded Gender (label) and Education (ordinal).
- Scaled numeric features only (Age, YearsExperience).
- Best model by RMSE: {best_model}
- Metrics:
{results_df.to_string()}
""")
```



```
Loaded dataset from: Salary_Data.csv
   Age Gender Education Level      Job Title  Years of Experience \
0  32.0   Male    Bachelor's  Software Engineer        5.0
1  28.0  Female    Master's   Data Analyst         3.0
2  45.0   Male        PhD  Senior Manager        15.0
3  36.0 Female    Bachelor's  Sales Associate        7.0
4  52.0   Male    Master's    Director          20.0
```

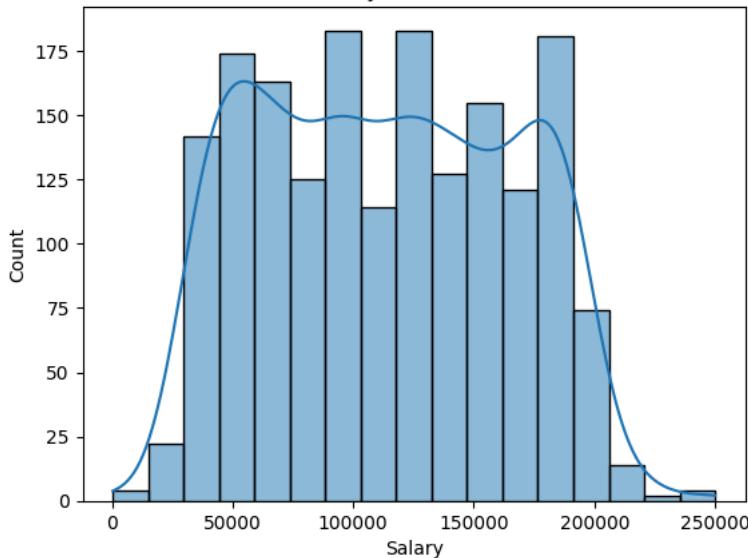
```
      Salary
0  90000.0
1  65000.0
2 150000.0
3  60000.0
4 200000.0
```

Descriptive statistics:

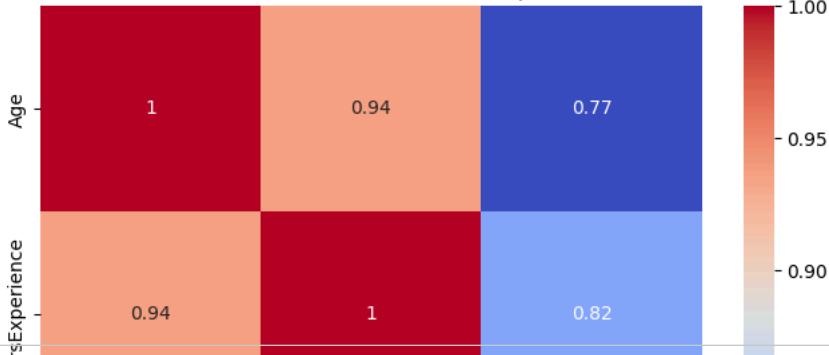
	Age	Gender	Education	JobTitle	\
count	1788.000000	1788.0	1788.0	1788	
unique	NaN	3.0	3.0	191	
top	NaN	1.0	5.0	Software Engineer	Manager
freq	NaN	967.0	1337.0		127
mean	35.135347	NaN	NaN		NaN
std	8.213003	NaN	NaN		NaN
min	21.000000	NaN	NaN		NaN
25%	29.000000	NaN	NaN		NaN
50%	33.000000	NaN	NaN		NaN
75%	41.000000	NaN	NaN		NaN
max	62.000000	NaN	NaN		NaN

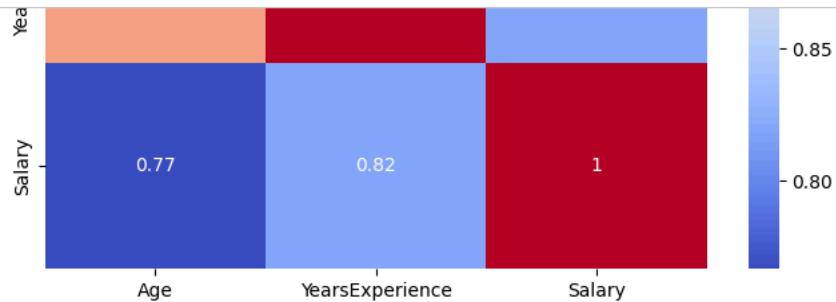
	YearsExperience	Salary
count	1788.000000	1788.000000
unique	NaN	NaN
top	NaN	NaN
freq	NaN	NaN
mean	9.154922	113177.285794
std	6.843199	51583.040514
min	0.000000	350.000000
25%	3.000000	70000.000000
50%	8.000000	110000.000000
75%	13.000000	160000.000000
max	34.000000	250000.000000

Salary Distribution



Feature Correlation Heatmap





Model performance:

	RMSE	R <sup>2</sup>
Random Forest	24898.261720	0.773298
Linear Regression	30918.765678	0.650408

Feature Importance (Random Forest)

